

**NoSQL Databases**

**Data Modeling**

**Aggregates + Graphs**

Rodrygo L. T. Santos

rodrygo@dcc.ufmg.br

# About data models

*“A data model [is] a collection of concepts that can be used to describe the structure of a database.”*

[Elmasri & Navathe, 2011]

*data types + relationships + constraints  
+ basic operations*

# The good ol' relational model

*users*

*attribute (column)*

*relation (table)*

<i>uid</i>	<i>firstname</i>	<i>lastname</i>
001	Joaquim	Carvalho
002	Manoel	Oliveira
003	João	Pinheiro

*tuple (row)*

*entity integrity constraint (primary key)*


# The good ol' relational model

users

<i>uid</i>	<i>firstname</i>	<i>lastname</i>
001	Joaquim	Carvalho
002	Manoel	Oliveira
003	João	Pinheiro

telephones

<i>uid</i>	<i>number</i>	<i>type</i>
001	3409-1234	work
001	8888-1234	mobile
002	3409-4321	work



*referential integrity constraint (foreign key)*

# The good ol' relational model

**SELECT**

    firstname, lastname, number, type

**FROM**

    users, telephones

**WHERE**

    users.uid = telephones.uid AND  
    lastname = 'Oliveira'

# Data modeling goals

- Data-oriented modeling
  - Driven by the structure of the data


*“What answers do I have?”*

- Query-oriented modeling
  - Driven by access patterns


*“What questions do I have?”*

# How to take advantage?

- **Principle of locality:** *“programs tend to use data and instructions with addresses **near or equal** to those they have used recently”*



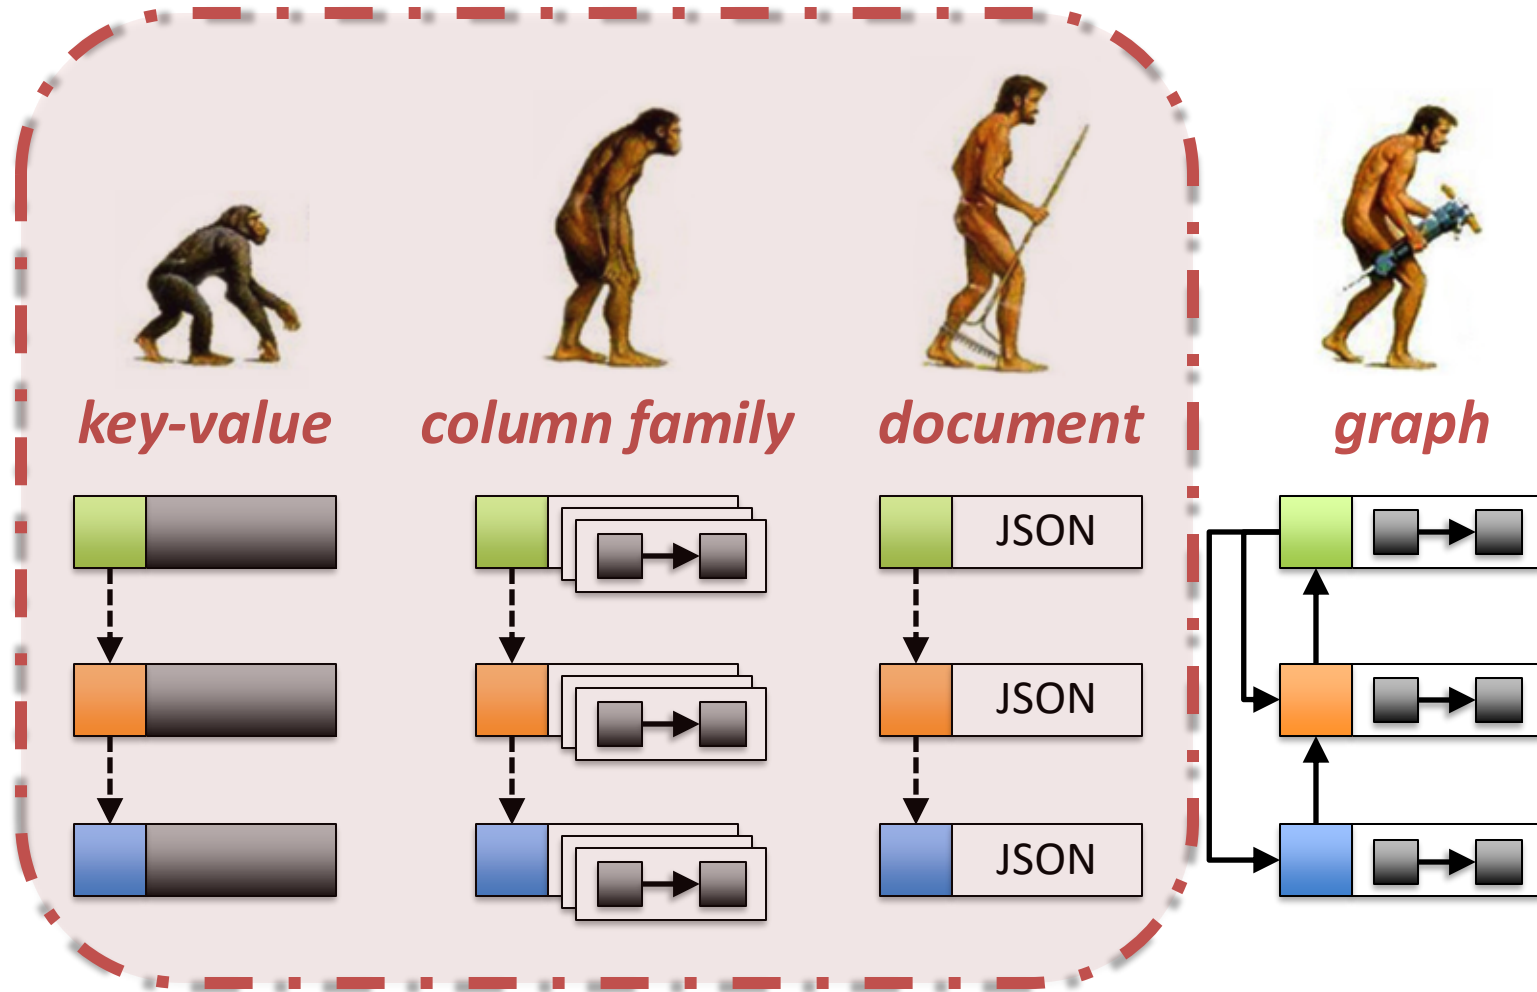
*temporal locality*



*spatial locality*

# NoSQL medicine

## *aggregate models*





# Aggregates

*“An aggregate is a collection of related objects that we wish to treat as a unit”*

[Fowler, 2013]

*orders*


*order lines*


*customers*


*credit cards*


*order*

ID: 1001

Customer: John Smith

Line items:

032129235 2 \$48 \$96

031986378 1 \$39 \$39

013765339 1 \$51 \$51

Payment details:

VISA \*\*\*\*\* 123

Expiry: 04/2020


# Why aggregates?

- Aggregates are not a logical data property
  - Remember query-driven modeling?
- Help greatly with running on a cluster
  - Minimize the number of nodes queried for data
- Help transaction processing
  - An aggregate is a natural transaction boundary
  - *Cross-aggregate atomicity needs application code*

# So, what flavors are there?

## Opaque aggregates

- Aggregate is a blob
  - Store whatever you like



- Can only query by key
  - Retrieve the whole thing

## Transparent aggregates

- Aggregate has structure
  - Must respect the underlying data structure (JSON, XML)



- Can query by key or value
  - Can retrieve parts or whole

# Key-value API

- Operations
  - get(key)
  - set(key, value, [ttl])
  - del(key)
- Data types
  - blob
  - list, set, hash of blob

# Key-value API

- Example (Riak)

```
curl -v -X POST -d '{ "lastVisit":1324669989288,
  "user":{"customerId":"91cfd5bcb7c",
  "name":"buyer",
  "countryCode":"US",
  "tzOffset":0}
}' \
-H "Content-Type: application/json" \
http://localhost:8098/buckets/session/keys/a7e618d9db25

curl -i http://localhost:8098/buckets/session/keys/a7e618d9db25
```

# Key-value DOs

- Session data
  - Session ID as the key
- Shopping cart data
  - Session ID as the key
- User profile data
  - User ID as the key
- ... and any other scenario that requires concurrently looking up things by key

# Column family model

	<i><b>“profile” column family</b></i>		<i><b>“orders” column family</b></i>		
<i><b>row key</b></i>	<i><b>column key</b></i>	<i><b>column key</b></i>	<i><b>column key</b></i>	<i><b>column key</b></i>	<i><b>column key</b></i>
1234	name	address	ORD1001	ORD1002	ORD1003
	t1 Martin	t1 5 <sup>th</sup> Ave.	t1 data	t1 data	t1 data
		t2 Broadway			
1235	name		ORD5001	ORD5002	
	t1 Pramod		t1 data	t1 data	

*row is the aggregate!*

# Column family API

- Operations
  - get(cfamily, row, [column])
  - set(cfamily, row, column, value)
  - del(cfamily, row, [column])  
or
  - CQL: a simple SQL flavor (Cassandra)
    - No joins, no subqueries, simple selection conditions
- Data types
  - blob



# Column family API

- Example (Cassandra CQL)

```
CREATE COLUMNFAMILY Customer (  
  KEY varchar PRIMARY KEY,  
  name varchar,  
  city varchar,  
  web varchar);
```

```
INSERT INTO Customer (KEY, name, city, web)  
VALUES ('mfowler', 'Martin Fowler', 'Boston',  
      'www.martinfowler.com');
```

```
SELECT name, web FROM Customer WHERE city='Boston'
```

# Column family DOs

- Event logging
  - Scalable writes
- Content management systems
  - Tags, categories, links, trackbacks as columns
- Counters
  - `INCR Visits['mfowler']['home'] BY 1;`
- Expiring data
  - `SET Customer['mfowler']['demo_access'] = 'allowed' WITH ttl=2592000;`

# Document API

- Operations
  - *get(XPath-like expression)*
  - *put(document)*
  - *del(docid)*
- Data types
  - Structured document (JSON, BSON, XML)

# Document API

- Example (MongoDB)

```
{
  "userId":"883c2c5",
  "firstname": "Pramod",
  "citiesvisited": [ "Chicago", "London" ],
  "addresses": [
    { "state": "AK",
      "city": "DILLINGHAM",
      "type": "R"
    },
    { "state": "MH",
      "city": "PUNE",
      "type": "R" }
  ],
  "lastcity": "London"
}
```

```
db.users.insert(document)
db.users.find({"userId":"883c2c5"})
db.users.find({"addresses.state":"/AK/})
```

# Document DOs

- Content management systems
  - Values are documents
- Search engines
  - Elasticsearch is a prominent example
- Data analytics
  - New bits can be added efficiently
- E-commerce applications
  - Flexible schema and cheap refactoring

# Aggregates

*“An aggregate is a collection of related objects that we wish to treat as a unit”*

[Fowler, 2013]

*orders*


*order lines*


*customers*


*credit cards*


*order*

ID: 1001

Customer: John Smith

Line items:

032129235 2 \$48 \$96

031986378 1 \$39 \$39

013765339 1 \$51 \$51

Payment details:

VISA \*\*\*\*\* 123

Expiry: 04/2020

# Aggregates

- Process customer together with an order
  - Order is the aggregate
- *Process each customer individually?*
  - Customer and order are the aggregates
  - Provide links between customer and order
  - Links can even be indexed
    - Link in transparent document fields
    - Link in metadata for opaque values

# WAIT!

*Customer and order in separate structures...  
Implicit links between the two...*

*This is not NoSQL...  
... this is OldSQL!*



# OldSQL to the rescue?

- Relationships in the relational world
  - Customer(customer\_id, customer\_name)  
Order(order\_id, ..., customer\_id)  
Order(customer\_id) **REFERENCES** Customer(customer\_id)
- Pull out your join hammer
  - SELECT \*  
FROM Customer **NATURAL JOIN** Order

# Writability

- Complex relationships = long SQL expressions
  - Customers who bought same product?  
**SELECT** C1.customer\_id, C2.customer\_id  
**FROM** Customer C1 **NATURAL JOIN** Order O1  
**NATURAL JOIN** Product P1 **NATURAL JOIN** Order O2  
**NATURAL JOIN** Customer C2  
**WHERE** C1.customer\_id <> C2.customer\_id
  - Customer and FOAF who bought same product?
    - Good luck...

# Efficiency

- Complex relationships = many index lookups
  - Path exists between two vertices?  
*(each person with an average of 50 friends)*


# persons	OldSQL
1,000	2,000 ms
1,000,000	?

# Enter graph databases

- OldSQL isn't cluster-friendly
  - Bring in aggregate databases
- OldSQL struggles with complex relationships
  - Bring in graph databases

# Remember graphs?

- The most flexible data structure



- But I could model a graph in OldSQL...  
... and in aggregate NoSQL databases as well

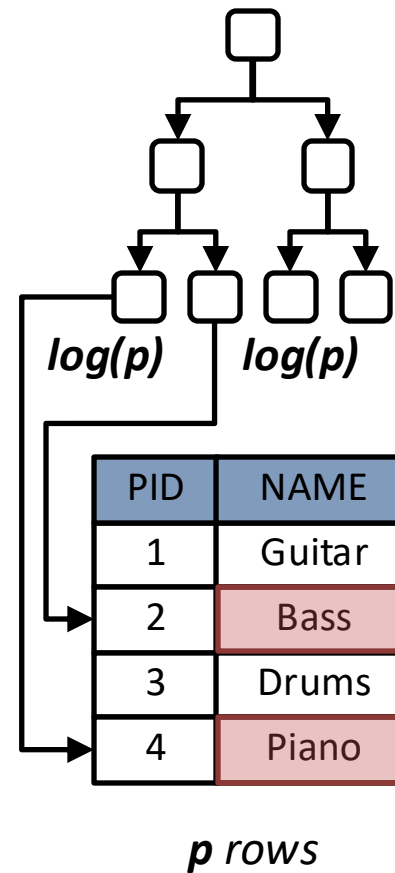
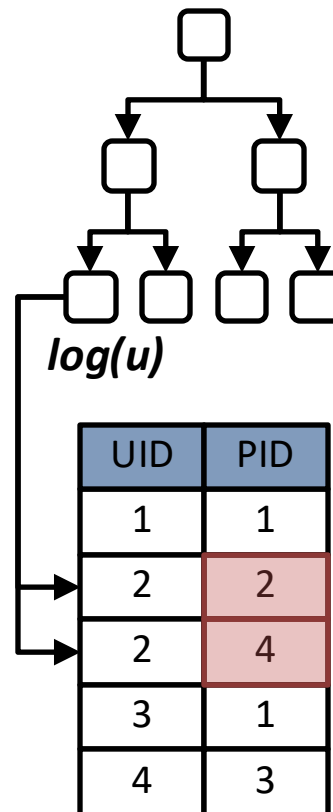

# What is a graph database?

*"A graph database [management system] is any storage system data provides **index-free** adjacencies."*

[Rodriguez and Neubauer, CoRR 2010]

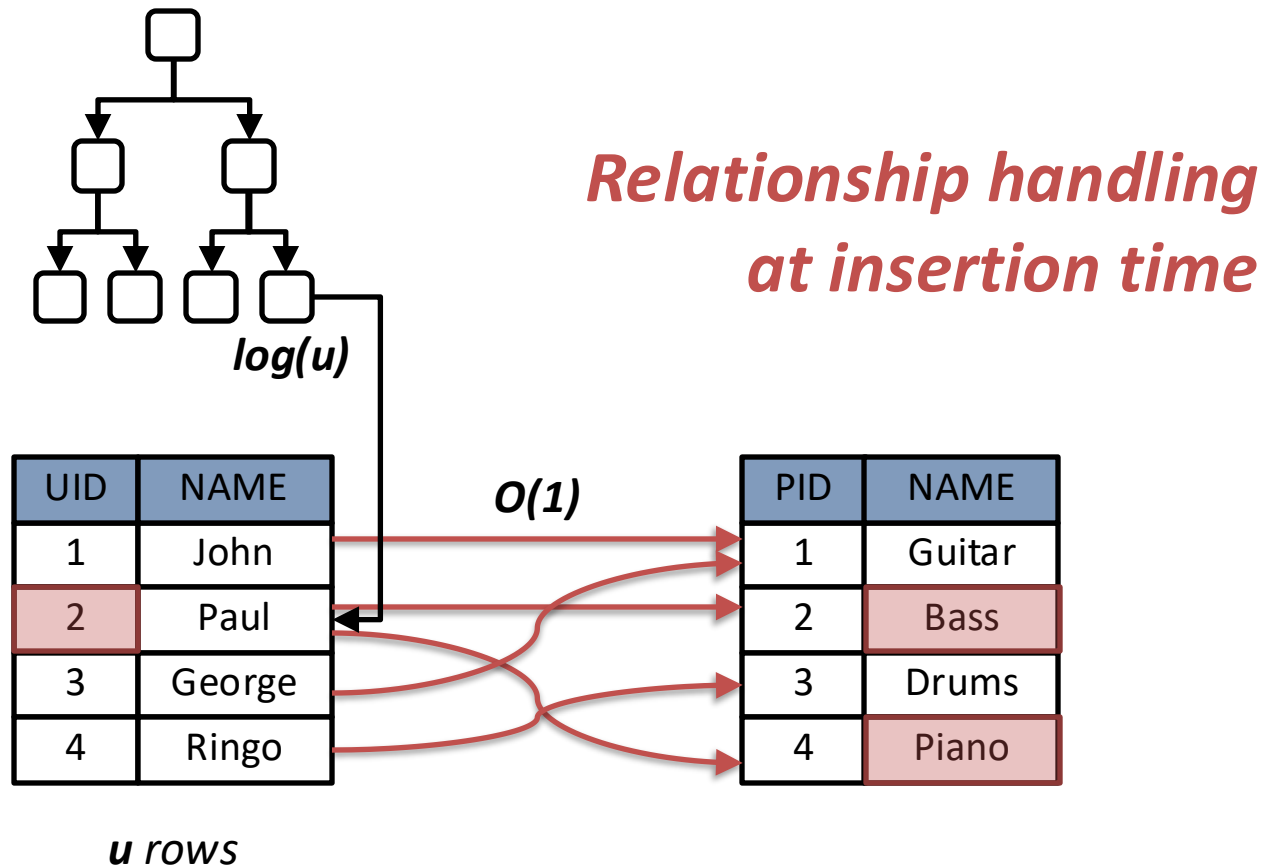
# Relational traversals

- What instruments does Paul have?



# Graph traversals

- What instruments does Paul have?



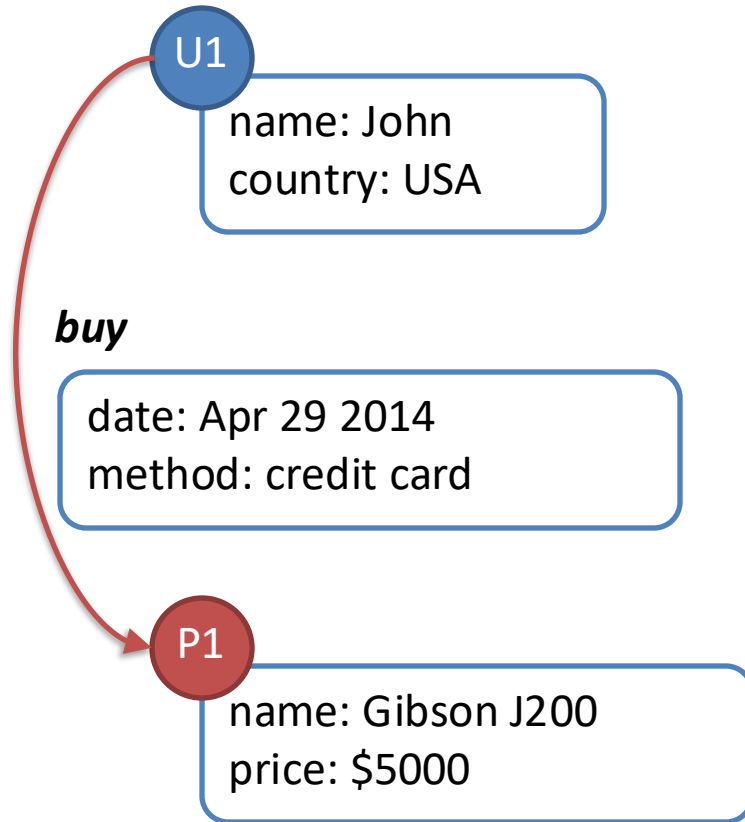


# Why graph databases?

- Complex relationships = many index lookups
  - Path exists between two vertices?  
*(each person with an average of 50 friends)*

# persons	OldSQL	Graph DB
1,000	2,000 ms	2 ms
1,000,000	Way too long...	2 ms

# The property graph model



- Directed, edge-labeled, attributed graph
  - Aka the property graph
- $G = (V, E, \lambda, \mu)$ 
  - $V$ : set of vertices
  - $E$ :  $E \subseteq (V \times V)$
  - $\lambda$ :  $E \rightarrow \Sigma$
  - $\Sigma$ : set of labels
  - $\mu$ :  $(V \cup E) \times R \rightarrow S$
  - $R$ : set of keys
  - $S$ : set of values

# Graph API

- Create, Update, Delete operations
  - add/del(vertex)
  - add/del(edge, vertex-in, vertex-out)
  - set/del(vertex/edge, key, value)
- Retrieve operation
  - Look up start vertex ( $O(\log n)$ )
  - Traverse edges ( $O(1)$ )

# Graph API

- Cypher (declarative)
  - Bob's friends and their location

```
START b = node:nodeIndex(name = "bob")
MATCH (b)-[:FRIEND]->(f)
RETURN f.name, f.location
```
  - Bob's friends' friends' friends

```
START b = node:nodeIndex(name = "bob")
MATCH path = (b)-[:FRIEND*1..3]->(f)
RETURN b.name, f.name, length(path)
```

# Graph API

- Gremlin (imperative)
  - Collaborative filtering for vertex 1

```
m = [:];  
g.v(1).out('likes').in('likes')  
  .out('likes').groupCount(m);  
m.sort{-it.value};
```
  - Vendor-agnostic standard
  - Extremely powerful
    - Transform, filter, side effect, branch predicates

# Graph DO's

- Connected data
  - Multi-relational graphs
- Location-based services
  - Place is a vertex
  - Edges have a distance property
- Recommendation engines
  - User and item as vertices
  - Click, rate, like, buy as edges